

In [1]:

Downloading the packages

```
!pip install xgboost
!pip install CatBoost
!pip install scikit-plot
!pip install autocorrect
```

In [2]:

Importing the necessary modules

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.utils import resample
from sklearn.ensemble import StackingClassifier
```

In [3]:

Load the reviews dataset

```
reviews = pd.read_csv('Amazon reviews.csv')
reviews.iloc[50:76]
```

Out[3]:

	category	rating	label	text
50	Home_and_Kitchen_5	5	0	Perfect. They do exactly what I need them to d...
51	Home_and_Kitchen_5	5	0	Love this movie & the time it took to finish. ...
52	Home_and_Kitchen_5	1	0	Will not seal properly. Atrocious little ones,...
53	Home_and_Kitchen_5	5	0	Got these for the third time. I have a small ...
54	Home_and_Kitchen_5	5	0	Excellent product and a much better quality th...
55	Home_and_Kitchen_5	5	1	These are just perfect, exactly what I was loo...
56	Home_and_Kitchen_5	5	1	Such a great purchase can't beat it for the price
57	Home_and_Kitchen_5	5	1	What can you say--- cheap and it works as inte...
58	Home_and_Kitchen_5	5	1	These are so nice, sturdy, like the color choi...
59	Home_and_Kitchen_5	5	1	It is nice bowl and have had a fast shipping!
60	Home_and_Kitchen_5	5	1	Great cup. Will last forever. Keeps things coo...
61	Home_and_Kitchen_5	5	1	Love them, just thought they would be a bit bi...
62	Home_and_Kitchen_5	5	1	Excellent quality product. Perfect for my ccoz...

63	Home_and_Kitchen_5	5	1	This fan is really pretty and I actually use it.
64	Home_and_Kitchen_5	1	1	Super rough, not soft wash cloths, more like b...
65	Home_and_Kitchen_5	5	1	Like this little guy. Use it often. He is small.
66	Home_and_Kitchen_5	5	1	Perfect size .easy to unfold. Topremoves for s...
67	Home_and_Kitchen_5	1	1	Does not do a very good job; difficult to use.
68	Home_and_Kitchen_5	5	1	frame was beautifully crafted and looks great ...
69	Home_and_Kitchen_5	5	1	My hot cocoa is finally perfect! This product ...
70	Home_and_Kitchen_5	5	1	i bought this for a friend and he loves it.
71	Home_and_Kitchen_5	5	1	Very soft and great quality for such a low price!
72	Home_and_Kitchen_5	5	1	Purchased as a Christmas gift. She will love it.
73	Home_and_Kitchen_5	5	1	accurate description, works great, came on tim...
74	Home_and_Kitchen_5	5	1	Use it in my Duvet cover. Fluffy and perfect fit!

In [4]:

```
# Text preprocessing
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
from nltk.tokenize import word_tokenize
from autocorrect import Speller

nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('punkt')

# Select a random row index
row_index = 2762

# Access the 'text' column before processing
print("Before processing:")
print(reviews['text'][row_index])
def preprocess_text(text):
    # Remove non-alphabetic characters
    text = re.sub('[^A-Za-z]', ' ', text)
    # Convert to lowercase
    text = text.lower()
    # Tokenize the text
    tokens = text.split()
    # Remove stopwords
    stopwords_list = stopwords.words('english')
```

```

    filtered_tokens = [token for token in tokens if token not in stopwords_list]
    # Stem the tokens
    stemmer = PorterStemmer()
    stemmed_tokens = [stemmer.stem(token) for token in filtered_tokens]
    # Lemmatize the tokens
    lemmatizer = WordNetLemmatizer()
    lemmatized_tokens = [lemmatizer.lemmatize(token) for token in stemmed_tokens]
    # Join the tokens back into a string
    preprocessed_text = ' '.join(lemmatized_tokens)
    return preprocessed_text

reviews['text'] = reviews['text'].apply(preprocess_text)

# Access the 'text' column after processing
print("After processing:")
print(reviews['text'][row_index])

```

Out [4]:

Before processing:

Very reliable. Originally purchased years ago for my apt, and when I moved, I bought a new one for a new home because (and this is my only complaint), the screws attaching it to the cabinet floor had rusted and stripped shut. The sliding mechanism, however, lasted years without oiling, and took tonnes of abuse and has never failed. Dead easy to install.

After processing:

reliabl origin purchas year ago apt move bought new one new home
complaint screw attach cabinet floor rust strip shut slide mechan
howev last year without oil took tonn abus never fail dead easi instal

In [5]:

```

# Separate the target variable and the text data
X = reviews['text']
y = reviews['label']

```

In [6]:

```

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
    test_size=0.3, random_state=42)

```

In [7]:

```
# Resample the training data to address data imbalance
X_train_resampled = []
y_train_resampled = []
for label in np.unique(y_train):
    X_train_label = X_train[y_train == label]
    y_train_label = y_train[y_train == label]
    X_train_label_resampled, y_train_label_resampled =
resample(X_train_label, y_train_label, n_samples=1000, replace=True,
random_state=42)
    X_train_resampled.append(X_train_label_resampled)
    y_train_resampled.append(y_train_label_resampled)
X_train_resampled = pd.concat(X_train_resampled)
y_train_resampled = pd.concat(y_train_resampled)
```

In [8]:

```
# Apply Tfidf vectorization to the text data
vectorizer = TfidfVectorizer(stop_words='english')
X_train_tfidf = vectorizer.fit_transform(X_train_resampled)
X_test_tfidf = vectorizer.transform(X_test)
```

In [9]:

```
# Apply standard scaling to the training and testing sets
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_tfidf.toarray())
X_test_scaled = scaler.transform(X_test_tfidf.toarray())
```

In [10]:

```
# Define the parameter grids for grid search
rf_param_grid = {'n_estimators': [100, 500, 1000], 'max_depth': [3, 5,
None]}
xgb_param_grid = {'n_estimators': [100, 500, 1000], 'max_depth': [3,
5, 7]}
lgb_param_grid = {'n_estimators': [100, 500, 1000], 'max_depth': [3,
5, 7]}
cb_param_grid = {'n_estimators': [100, 500, 1000], 'max_depth': [3, 5,
7]}
gb_param_grid = {'n_estimators': [100, 500, 1000], 'max_depth': [3, 5,
7]}
```

In [11]:

```
# Perform grid search to optimize the classifiers
rf_grid = GridSearchCV(RandomForestClassifier(random_state=42),
rf_param_grid, cv=5, n_jobs=-1)
xgb_grid = GridSearchCV(XGBClassifier(random_state=42),
xgb_param_grid, cv=5, n_jobs=-1)
```

```

lgb_grid = GridSearchCV(LGBMClassifier(random_state=42),
lgb_param_grid, cv=5, n_jobs=-1)
cb_grid = GridSearchCV(CatBoostClassifier(random_state=42),
cb_param_grid, cv=5, n_jobs=-1)
gb_grid = GridSearchCV(RandomForestClassifier(random_state=42),
gb_param_grid, cv=5, n_jobs=-1)

```

In [12]:

Fit the optimized classifiers to the training data

```

rf_grid.fit(X_train_scaled, y_train_resampled)
xgb_grid.fit(X_train_scaled, y_train_resampled)
lgb_grid.fit(X_train_scaled, y_train_resampled)
cb_grid.fit(X_train_scaled, y_train_resampled)
gb_grid.fit(X_train_scaled, y_train_resampled)

```

Out [12]:

Learning rate set to 0.013851

0:	learn: 0.6909365 total: 166ms	remaining: 2m 45s
1:	learn: 0.6891085 total: 256ms	remaining: 2m 7s
2:	learn: 0.6869619 total: 373ms	remaining: 2m 3s
3:	learn: 0.6849605 total: 470ms	remaining: 1m 56s
4:	learn: 0.6828380 total: 562ms	remaining: 1m 51s
5:	learn: 0.6810862 total: 654ms	remaining: 1m 48s
6:	learn: 0.6791836 total: 742ms	remaining: 1m 45s
7:	learn: 0.6771038 total: 831ms	remaining: 1m 43s
8:	learn: 0.6748765 total: 925ms	remaining: 1m 41s
9:	learn: 0.6733539 total: 1.01s	remaining: 1m 40s
10:	learn: 0.6715230 total: 1.1s	remaining: 1m 38s
11:	learn: 0.6698818 total: 1.19s	remaining: 1m 37s
12:	learn: 0.6682977 total: 1.28s	remaining: 1m 36s
13:	learn: 0.6668198 total: 1.37s	remaining: 1m 36s
14:	learn: 0.6651077 total: 1.47s	remaining: 1m 36s
15:	learn: 0.6636890 total: 1.56s	remaining: 1m 35s

```

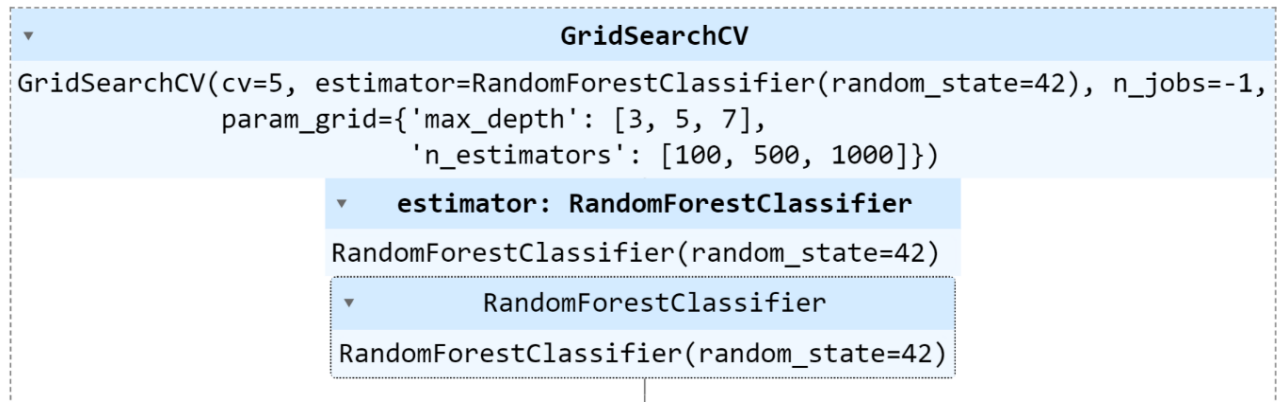
.....
981: learn: 0.2974422 total: 1m 37s    remaining: 1.79s
982: learn: 0.2971875 total: 1m 37s    remaining: 1.69s
983: learn: 0.2969678 total: 1m 37s    remaining: 1.59s
984: learn: 0.2966740 total: 1m 38s    remaining: 1.49s
985: learn: 0.2964466 total: 1m 38s    remaining: 1.39s
986: learn: 0.2961845 total: 1m 38s    remaining: 1.29s
987: learn: 0.2960369 total: 1m 38s    remaining: 1.19s
988: learn: 0.2957993 total: 1m 38s    remaining: 1.09s
989: learn: 0.2956468 total: 1m 38s    remaining: 996ms
990: learn: 0.2954216 total: 1m 38s    remaining: 897ms
991: learn: 0.2951882 total: 1m 38s    remaining: 798ms
992: learn: 0.2949301 total: 1m 39s    remaining: 698ms

```

```

993: learn: 0.2947020 total: 1m 39s    remaining: 599ms
994: learn: 0.2945699 total: 1m 39s    remaining: 499ms
995: learn: 0.2944823 total: 1m 39s    remaining: 400ms
996: learn: 0.2943984 total: 1m 39s    remaining: 300ms
997: learn: 0.2942040 total: 1m 39s    remaining: 200ms
998: learn: 0.2941395 total: 1m 40s    remaining: 100ms
999: learn: 0.2940769 total: 1m 40s    remaining: 0us

```



In [13]:

Extract the best classifiers from the grid search results

```

rf_best = rf_grid.best_estimator_
xgb_best = xgb_grid.best_estimator_
lgb_best = lgb_grid.best_estimator_
cb_best = cb_grid.best_estimator_
gb_best = gb_grid.best_estimator_

```

In [14]:

Define the ensemble classifiers using majority voting and stacking

```

voting_clf = VotingClassifier([('rf', rf_best), ('xgb', xgb_best),
                                ('lgb', lgb_best), ('cb', cb_best), ('gb', gb_best)], voting='soft')
stacking_clf = StackingClassifier([('rf', rf_best), ('xgb', xgb_best),
                                     ('lgb', lgb_best), ('cb', cb_best), ('gb', gb_best)])

```

In [15]:

Fit the ensemble classifiers to the training data

```

voting_clf.fit(X_train_scaled, y_train_resampled)
stacking_clf.fit(X_train_scaled, y_train_resampled)

```

Out [15]:

Learning rate set to 0.012592

```

4: learn: 0.6846217    total: 375ms    remaining: 1m 14s
5: learn: 0.6830193    total: 463ms    remaining: 1m 16s
6: learn: 0.6817952    total: 532ms    remaining: 1m 15s
7: learn: 0.6799940    total: 600ms    remaining: 1m 14s
8: learn: 0.6783643    total: 673ms    remaining: 1m 14s

```

9:	learn: 0.6769998	total: 749ms	remaining: 1m 14s
10:	learn: 0.6756055	total: 820ms	remaining: 1m 13s
11:	learn: 0.6744568	total: 894ms	remaining: 1m 13s
12:	learn: 0.6730493	total: 965ms	remaining: 1m 13s
13:	learn: 0.6712112	total: 1.04s	remaining: 1m 13s
14:	learn: 0.6694191	total: 1.11s	remaining: 1m 13s
15:	learn: 0.6681017	total: 1.19s	remaining: 1m 12s

.....

981:	learn: 0.3078034	total: 1m 20s	remaining: 1.47s
982:	learn: 0.3076445	total: 1m 20s	remaining: 1.39s
983:	learn: 0.3075808	total: 1m 20s	remaining: 1.3s
984:	learn: 0.3073479	total: 1m 20s	remaining: 1.22s
985:	learn: 0.3071265	total: 1m 20s	remaining: 1.14s
986:	learn: 0.3070223	total: 1m 20s	remaining: 1.06s
987:	learn: 0.3068619	total: 1m 20s	remaining: 977ms
988:	learn: 0.3066641	total: 1m 20s	remaining: 896ms
989:	learn: 0.3065057	total: 1m 20s	remaining: 814ms
990:	learn: 0.3063939	total: 1m 20s	remaining: 733ms
991:	learn: 0.3061856	total: 1m 20s	remaining: 651ms
992:	learn: 0.3060314	total: 1m 20s	remaining: 570ms
993:	learn: 0.3056331	total: 1m 20s	remaining: 488ms
994:	learn: 0.3053640	total: 1m 20s	remaining: 407ms
995:	learn: 0.3051338	total: 1m 21s	remaining: 326ms
996:	learn: 0.3050367	total: 1m 21s	remaining: 244ms
997:	learn: 0.3048031	total: 1m 21s	remaining: 163ms
998:	learn: 0.3045684	total: 1m 21s	remaining: 81.4ms
999:	learn: 0.3044225	total: 1m 21s	remaining: 0us

Learning rate set to 0.012592

0:	learn: 0.6908525	total: 88.5ms	remaining: 1m 28s
1:	learn: 0.6886115	total: 158ms	remaining: 1m 18s
2:	learn: 0.6870465	total: 233ms	remaining: 1m 17s
3:	learn: 0.6852878	total: 307ms	remaining: 1m 16s
4:	learn: 0.6840765	total: 375ms	remaining: 1m 14s
5:	learn: 0.6821594	total: 458ms	remaining: 1m 15s
6:	learn: 0.6804935	total: 534ms	remaining: 1m 15s
7:	learn: 0.6790381	total: 613ms	remaining: 1m 15s
8:	learn: 0.6773534	total: 694ms	remaining: 1m 16s
9:	learn: 0.6758022	total: 764ms	remaining: 1m 15s
10:	learn: 0.6744437	total: 833ms	remaining: 1m 14s
11:	learn: 0.6733758	total: 904ms	remaining: 1m 14s
12:	learn: 0.6717796	total: 976ms	remaining: 1m 14s
13:	learn: 0.6703937	total: 1.05s	remaining: 1m 13s
14:	learn: 0.6690930	total: 1.13s	remaining: 1m 13s
15:	learn: 0.6671332	total: 1.2s	remaining: 1m 13s

.....

981:	learn: 0.3023393	total: 1m 20s	remaining: 1.47s
982:	learn: 0.3021164	total: 1m 20s	remaining: 1.39s
983:	learn: 0.3018658	total: 1m 20s	remaining: 1.31s
984:	learn: 0.3016576	total: 1m 20s	remaining: 1.23s
985:	learn: 0.3014143	total: 1m 20s	remaining: 1.15s
986:	learn: 0.3011494	total: 1m 20s	remaining: 1.07s
987:	learn: 0.3010367	total: 1m 21s	remaining: 985ms
988:	learn: 0.3008613	total: 1m 21s	remaining: 903ms
989:	learn: 0.3007253	total: 1m 21s	remaining: 822ms
990:	learn: 0.3004653	total: 1m 21s	remaining: 740ms
991:	learn: 0.3001637	total: 1m 21s	remaining: 658ms
992:	learn: 0.2999900	total: 1m 21s	remaining: 576ms
993:	learn: 0.2999283	total: 1m 21s	remaining: 494ms
994:	learn: 0.2997671	total: 1m 21s	remaining: 412ms
995:	learn: 0.2997016	total: 1m 21s	remaining: 329ms
996:	learn: 0.2995307	total: 1m 22s	remaining: 247ms
997:	learn: 0.2993978	total: 1m 22s	remaining: 165ms
998:	learn: 0.2991948	total: 1m 22s	remaining: 82.3ms
999:	learn: 0.2989196	total: 1m 22s	remaining: 0us

Learning rate set to 0.012592

0:	learn: 0.6920444	total: 88.2ms	remaining: 1m 28s
1:	learn: 0.6897713	total: 162ms	remaining: 1m 20s
2:	learn: 0.6878114	total: 239ms	remaining: 1m 19s
3:	learn: 0.6859438	total: 311ms	remaining: 1m 17s
4:	learn: 0.6846315	total: 383ms	remaining: 1m 16s
5:	learn: 0.6829251	total: 467ms	remaining: 1m 17s
6:	learn: 0.6806837	total: 538ms	remaining: 1m 16s
7:	learn: 0.6790530	total: 611ms	remaining: 1m 15s
8:	learn: 0.6768527	total: 698ms	remaining: 1m 16s
9:	learn: 0.6751366	total: 772ms	remaining: 1m 16s
10:	learn: 0.6730441	total: 851ms	remaining: 1m 16s
11:	learn: 0.6711665	total: 931ms	remaining: 1m 16s
12:	learn: 0.6701586	total: 1s	remaining: 1m 16s
13:	learn: 0.6687129	total: 1.08s	remaining: 1m 15s
14:	learn: 0.6669098	total: 1.16s	remaining: 1m 15s
15:	learn: 0.6655400	total: 1.23s	remaining: 1m 15s

.....

981:	learn: 0.3037858	total: 1m 22s	remaining: 1.5s
982:	learn: 0.3035218	total: 1m 22s	remaining: 1.42s
983:	learn: 0.3033497	total: 1m 22s	remaining: 1.34s
984:	learn: 0.3030268	total: 1m 22s	remaining: 1.25s

985:	learn: 0.3027258	total: 1m 22s	remaining: 1.17s
986:	learn: 0.3025821	total: 1m 22s	remaining: 1.09s
987:	learn: 0.3024016	total: 1m 22s	remaining: 1s
988:	learn: 0.3020612	total: 1m 22s	remaining: 919ms
989:	learn: 0.3019968	total: 1m 22s	remaining: 836ms
990:	learn: 0.3019332	total: 1m 22s	remaining: 752ms
991:	learn: 0.3017925	total: 1m 22s	remaining: 668ms
992:	learn: 0.3014855	total: 1m 22s	remaining: 585ms
993:	learn: 0.3012717	total: 1m 23s	remaining: 501ms
994:	learn: 0.3010373	total: 1m 23s	remaining: 418ms
995:	learn: 0.3008316	total: 1m 23s	remaining: 334ms
996:	learn: 0.3005971	total: 1m 23s	remaining: 251ms
997:	learn: 0.3003900	total: 1m 23s	remaining: 167ms
998:	learn: 0.3002767	total: 1m 23s	remaining: 83.5ms
999:	learn: 0.3001432	total: 1m 23s	remaining: 0us

Learning rate set to 0.012592

0:	learn: 0.6916444	total: 87.7ms	remaining: 1m 27s
1:	learn: 0.6892218	total: 160ms	remaining: 1m 19s
2:	learn: 0.6866895	total: 239ms	remaining: 1m 19s
3:	learn: 0.6849472	total: 310ms	remaining: 1m 17s
4:	learn: 0.6832336	total: 381ms	remaining: 1m 15s
5:	learn: 0.6814099	total: 464ms	remaining: 1m 16s
6:	learn: 0.6795859	total: 536ms	remaining: 1m 16s
7:	learn: 0.6779248	total: 611ms	remaining: 1m 15s
8:	learn: 0.6763097	total: 698ms	remaining: 1m 16s
9:	learn: 0.6745213	total: 769ms	remaining: 1m 16s
10:	learn: 0.6729563	total: 840ms	remaining: 1m 15s
11:	learn: 0.6715945	total: 919ms	remaining: 1m 15s
12:	learn: 0.6701184	total: 990ms	remaining: 1m 15s
13:	learn: 0.6680466	total: 1.06s	remaining: 1m 14s
14:	learn: 0.6666321	total: 1.16s	remaining: 1m 16s
15:	learn: 0.6653995	total: 1.23s	remaining: 1m 15s

.....

981:	learn: 0.3041048	total: 1m 21s	remaining: 1.49s
982:	learn: 0.3039346	total: 1m 21s	remaining: 1.41s
983:	learn: 0.3038754	total: 1m 21s	remaining: 1.33s
984:	learn: 0.3037389	total: 1m 21s	remaining: 1.24s
985:	learn: 0.3035018	total: 1m 21s	remaining: 1.16s
986:	learn: 0.3032521	total: 1m 21s	remaining: 1.08s
987:	learn: 0.3029439	total: 1m 22s	remaining: 997ms
988:	learn: 0.3027430	total: 1m 22s	remaining: 914ms

989:	learn: 0.3026220	total: 1m 22s	remaining: 832ms
990:	learn: 0.3024990	total: 1m 22s	remaining: 748ms
991:	learn: 0.3022544	total: 1m 22s	remaining: 665ms
992:	learn: 0.3021512	total: 1m 22s	remaining: 582ms
993:	learn: 0.3020348	total: 1m 22s	remaining: 499ms
994:	learn: 0.3018712	total: 1m 22s	remaining: 416ms
995:	learn: 0.3017743	total: 1m 22s	remaining: 332ms
996:	learn: 0.3016407	total: 1m 22s	remaining: 249ms
997:	learn: 0.3015196	total: 1m 22s	remaining: 166ms
998:	learn: 0.3014478	total: 1m 23s	remaining: 83.1ms
999:	learn: 0.3013050	total: 1m 23s	remaining: 0us

Learning rate set to 0.012592

0:	learn: 0.6911581	total: 92ms	remaining: 1m 31s
1:	learn: 0.6892107	total: 179ms	remaining: 1m 29s
2:	learn: 0.6873435	total: 252ms	remaining: 1m 23s
3:	learn: 0.6850827	total: 323ms	remaining: 1m 20s
4:	learn: 0.6833879	total: 401ms	remaining: 1m 19s
5:	learn: 0.6817052	total: 474ms	remaining: 1m 18s
6:	learn: 0.6801662	total: 543ms	remaining: 1m 17s
7:	learn: 0.6786194	total: 622ms	remaining: 1m 17s
8:	learn: 0.6769754	total: 695ms	remaining: 1m 16s
9:	learn: 0.6751120	total: 766ms	remaining: 1m 15s
10:	learn: 0.6733614	total: 843ms	remaining: 1m 15s
11:	learn: 0.6723055	total: 920ms	remaining: 1m 15s
12:	learn: 0.6706668	total: 992ms	remaining: 1m 15s
13:	learn: 0.6689574	total: 1.07s	remaining: 1m 15s
14:	learn: 0.6674405	total: 1.15s	remaining: 1m 15s
15:	learn: 0.6659634	total: 1.23s	remaining: 1m 15s

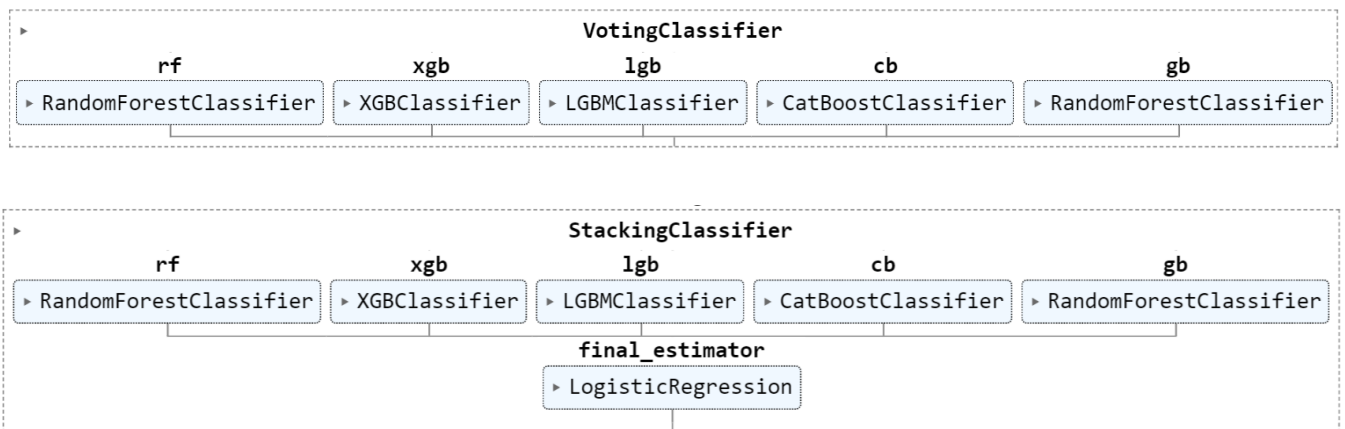
.....

981:	learn: 0.3011695	total: 1m 19s	remaining: 1.46s
982:	learn: 0.3009254	total: 1m 19s	remaining: 1.38s
983:	learn: 0.3007502	total: 1m 19s	remaining: 1.3s
984:	learn: 0.3006521	total: 1m 20s	remaining: 1.22s
985:	learn: 0.3004708	total: 1m 20s	remaining: 1.14s
986:	learn: 0.3003749	total: 1m 20s	remaining: 1.06s
987:	learn: 0.3001303	total: 1m 20s	remaining: 975ms
988:	learn: 0.2999314	total: 1m 20s	remaining: 894ms
989:	learn: 0.2996229	total: 1m 20s	remaining: 813ms
990:	learn: 0.2993063	total: 1m 20s	remaining: 732ms
991:	learn: 0.2991958	total: 1m 20s	remaining: 651ms
992:	learn: 0.2989733	total: 1m 20s	remaining: 570ms
993:	learn: 0.2987368	total: 1m 20s	remaining: 489ms

```

994:    learn: 0.2985793    total: 1m 21s    remaining: 408ms
995:    learn: 0.2983046    total: 1m 21s    remaining: 326ms
996:    learn: 0.2981743    total: 1m 21s    remaining: 245ms
997:    learn: 0.2980621    total: 1m 21s    remaining: 163ms
998:    learn: 0.2979987    total: 1m 21s    remaining: 81.7ms
999:    learn: 0.2979361    total: 1m 21s    remaining: 0us

```



In [16]:

```

# Make predictions using the ensemble classifiers
y_pred_voting = voting_clf.predict(X_test_scaled)
y_pred_stacking = stacking_clf.predict(X_test_scaled)

```

In [17]:

```

# Get the accuracy of the ensemble classifiers
voting_acc = accuracy_score(y_test, y_pred_voting)
stacking_acc = accuracy_score(y_test, y_pred_stacking)

```

In [18]:

```

from sklearn.metrics import accuracy_score
# Get the accuracy of each classifier
rf_acc = accuracy_score(y_test, rf_best.predict(X_test_scaled))
xgb_acc = accuracy_score(y_test, xgb_best.predict(X_test_scaled))
lgb_acc = accuracy_score(y_test, lgb_best.predict(X_test_scaled))
cb_acc = accuracy_score(y_test, cb_best.predict(X_test_scaled))
gb_acc = accuracy_score(y_test, gb_best.predict(X_test_scaled))

# Print the accuracies
print('Random Forest Accuracy: %.2f%%' % (rf_acc * 100))
print('XGBoost Accuracy: %.2f%%' % (xgb_acc * 100))
print('LightGBM Accuracy: %.2f%%' % (lgb_acc * 100))
print('CatBoost Accuracy: %.2f%%' % (cb_acc * 100))
print('Gradient Boosting Accuracy: %.2f%%' % (gb_acc * 100))

```

```
print("\n")
# Print the accuracy of the ensemble classifiers
print('Voting Classifier Accuracy: %.2f%%' % (voting_acc * 100))
print('Stacking Classifier Accuracy: %.2f%%' % (stacking_acc * 100))
```

Out [18]:

```
Random Forest Accuracy: 76.72%
XGBoost Accuracy: 75.75%
LightGBM Accuracy: 76.08%
CatBoost Accuracy: 77.00%
Gradient Boosting Accuracy: 74.87%
```

```
Voting Classifier Accuracy: 77.33%
Stacking Classifier Accuracy: 77.39%
```

In [19]:


```
import matplotlib.pyplot as plt
import numpy as np

models = ['Random Forest', 'XGBoost', 'Light GBM', 'Cat Boost',
          'Gradient Boosting', 'Voting', 'Stacking']
accuracies = [rf_acc, xgb_acc, lgb_acc, cb_acc, gb_acc, voting_acc,
              stacking_acc]

# create bar plot
plt.bar(models, accuracies, color=['blue', 'gold', 'green', 'cyan',
                                   'red', 'magenta', 'coral' ])
plt.xticks(np.arange(len(models)), models, rotation=45, ha='right')

plt.title('Accuracy Comparison')
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.ylim(0.72, 0.78)
plt.show()
```

Out [19]:



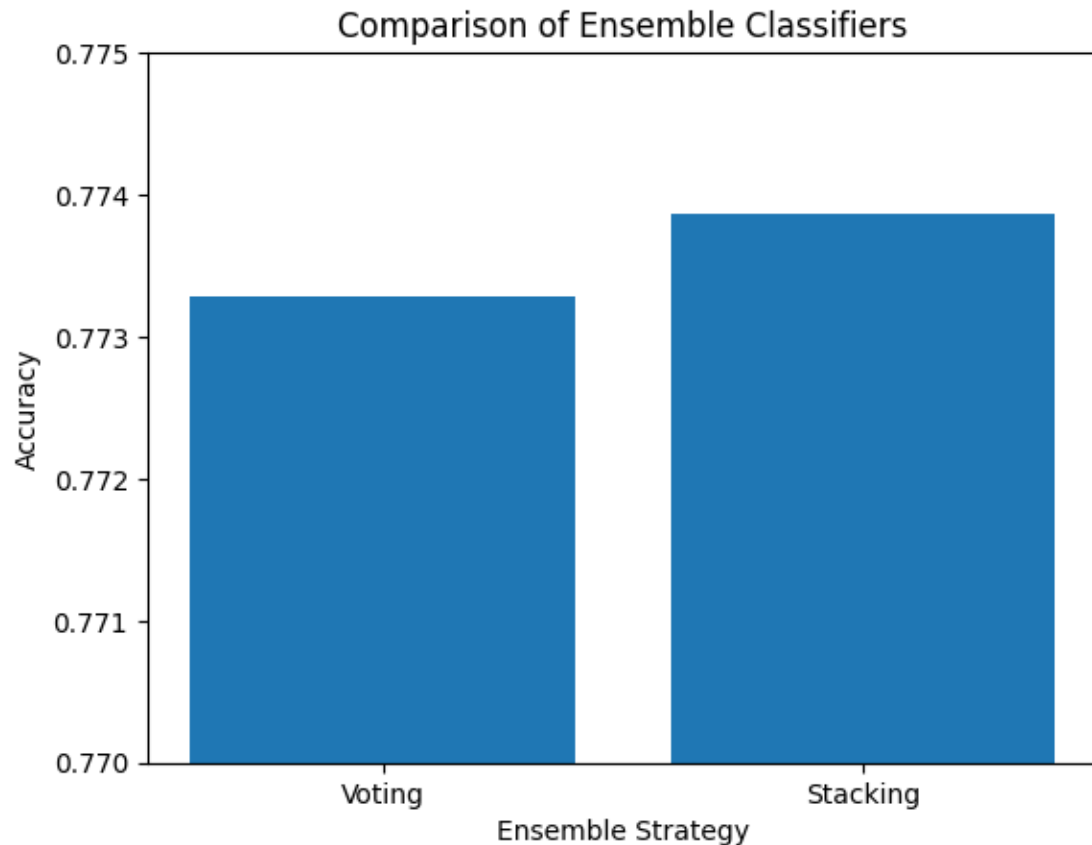
In [20]:

```
import numpy as np
import matplotlib.pyplot as plt

# Create a bar chart to compare the accuracies of the two ensemble
classifiers
models = ['Voting', 'Stacking']
accuracies = [voting_acc, stacking_acc]

x_pos = np.arange(len(models))
plt.bar(x_pos, accuracies)
plt.ylim(0.77, 0.775)
plt.title('Comparison of Ensemble Classifiers')
plt.xlabel('Ensemble Strategy')
plt.ylabel('Accuracy')
plt.xticks(x_pos, models)
plt.show()
```

Out [20]:



In [21]:

```
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import confusion_matrix

# Define a function to plot the confusion matrix
def plot_confusion_matrix(cm, classes, normalize=False,
                           cmap=plt.cm.Blues):
    """
    This function plots the confusion matrix.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title('Confusion Matrix')
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 2.0
```

```

for i, j in np.ndindex(cm.shape):
    plt.text(j, i, format(cm[i, j], '.2f' if normalize else 'd'),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

# Calculate the confusion matrix for the voting classifier
conf_matrix_voting = confusion_matrix(y_test, y_pred_voting)
print('Confusion Matrix (voting):\n', conf_matrix_voting)
plot_confusion_matrix(conf_matrix_voting, classes=['Negative',
'Positive'])
print("\n")
# Calculate the confusion matrix for the stacking classifier
conf_matrix_stacking = confusion_matrix(y_test, y_pred_stacking)
print('Confusion Matrix (stacking):\n', conf_matrix_stacking)
plot_confusion_matrix(conf_matrix_stacking, classes=['Negative',
'Positive'])
print("\n")

```

Out [21]:

Confusion Matrix (voting):

```

[[4702 1364]
 [1386 4678]]


```

Confusion Matrix (stacking):

```

[[4844 1222]
 [1521 4543]]

```



In [22]:

Get the TP, FP, TN, and FN rates for the voting classifier

```
tn_voting, fp_voting, fn_voting, tp_voting = conf_matrix_voting.ravel(
)
```

```
tpr_voting = tp_voting / (tp_voting + fn_voting)
```

```
fpr_voting = fp_voting / (fp_voting + tn_voting)
```

```
tnr_voting = tn_voting / (tn_voting + fp_voting)
```

```
fnr_voting = fn_voting / (tp_voting + fn_voting)
```

Print the rates for the voting classifier

```
print("TP rate (voting):", tpr_voting)
```

```
print("FP rate (voting):", fpr_voting)
```

```
print("TN rate (voting):", tnr_voting)
```

```
print("FN rate (voting):", fnr_voting)
```

Get the TP, FP, TN, and FN rates for the stacking classifier

```
tn_stacking, fp_stacking, fn_stacking, tp_stacking = conf_matrix_stack
ing.ravel()
```

```
tpr_stacking = tp_stacking / (tp_stacking + fn_stacking)
```

```
fpr_stacking = fp_stacking / (fp_stacking + tn_stacking)
```

```
tnr_stacking = tn_stacking / (tn_stacking + fp_stacking)
```

```
fnr_stacking = fn_stacking / (tp_stacking + fn_stacking)
```



```
# Print the rates for the stacking classifier
```

```
print("TP rate (stacking):", tpr_stacking)
print("FP rate (stacking):", fpr_stacking)
print("TN rate (stacking):", tnr_stacking)
print("FN rate (stacking):", fnr_stacking)
```

```
Out [22]:
```

```
TP rate (voting): 0.7714379947229552
FP rate (voting): 0.22485987471150676
TN rate (voting): 0.7751401252884932
FN rate (voting): 0.22856200527704484
TP rate (stacking): 0.7491754617414248
FP rate (stacking): 0.2014507088691065
TN rate (stacking): 0.7985492911308935
FN rate (stacking): 0.2508245382585752
```

```
In [23]:
```

```
import matplotlib.pyplot as plt
```

```
# Create lists of the rates for the voting and stacking classifiers
```

```
tpr_list = [tpr_voting, tpr_stacking]
fpr_list = [fpr_voting, fpr_stacking]
tnr_list = [tnr_voting, tnr_stacking]
fnr_list = [fnr_voting, fnr_stacking]
```

```
# Create a list of the classifier names
```

```
classifiers = ['Voting', 'Stacking']
```

```
# Create a bar plot
```

```
fig, ax = plt.subplots(figsize=(10,6))
ax.bar(classifiers, tpr_list, color='r', alpha=opacity, label='TP rate')
ax.bar(classifiers, fpr_list, color='b', alpha=opacity, label='FP rate',
        bottom=tpr_list)
ax.bar(classifiers, tnr_list, color='g', alpha=opacity, label='TN rate',
        bottom=[sum(x) for x in zip(tpr_list, fpr_list)])
ax.bar(classifiers, fnr_list, color='y', alpha=opacity, label='FN rate',
        bottom=[sum(x) for x in zip(tpr_list, fpr_list, tnr_list)])
```

```
# Add value labels to the bars
```

```
for i, v in enumerate(tpr_list):
    ax.text(i, v/2, "{:.2f}".format(v), color='white', fontweight='bold',
            ha='center', va='center')
    ax.text(i, v+fpr_list[i]/2, "{:.2f}".format(fpr_list[i]), color='white',
            fontweight='bold', ha='center', va='center')
```

```

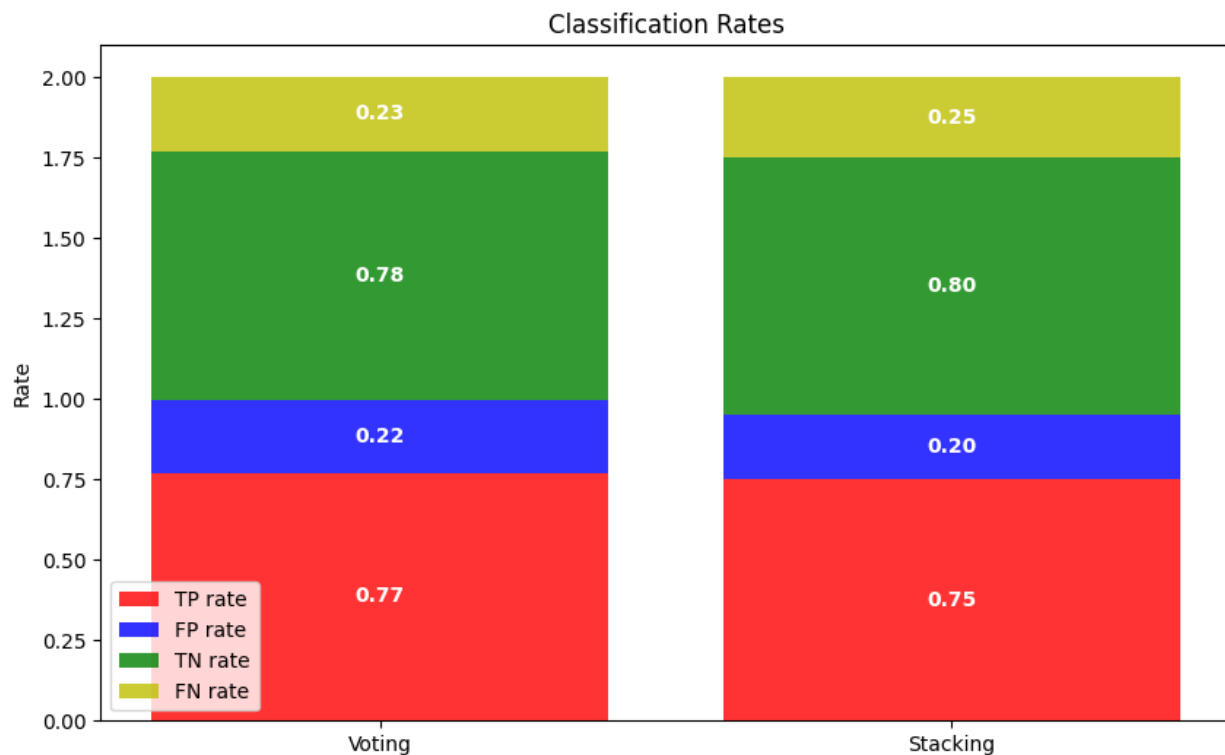
    ax.text(i, v+fpr_list[i]+tnr_list[i]/2, "{:.2f}".format(tnr_list[i]
)], color='white', fontweight='bold', ha='center', va='center')
    ax.text(i, v+fpr_list[i]+tnr_list[i]+fnr_list[i]/2, "{:.2f}".forma
t(fnr_list[i]), color='white', fontweight='bold', ha='center', va='cen
ter')

# Add labels and legend
ax.set_ylabel('Rate')
ax.set_title('Classification Rates')
ax.legend()

# Show the plot
plt.show()

```

Out [23]:



In [24]:

```

# Calculate classification report for voting & Stacking classifier
print('Classification Report (voting):\n',
classification_report(y_test, y_pred_voting))

print('Classification Report (stacking):\n',

```

```

classification_report(y_test, y_pred_stacking))

# Calculate F1 score, recall, precision, and accuracy for voting
classifier
voting_f1 = classification_report(y_test, y_pred_voting,
output_dict=True)['weighted avg']['f1-score']
voting_recall = classification_report(y_test, y_pred_voting,
output_dict=True)['weighted avg']['recall']
voting_precision = classification_report(y_test, y_pred_voting,
output_dict=True)['weighted avg']['precision']
voting_acc = accuracy_score(y_test, y_pred_voting)

# Calculate F1 score, recall, precision, and accuracy for stacking
classifier
stacking_f1 = classification_report(y_test, y_pred_stacking,
output_dict=True)['weighted avg']['f1-score']
stacking_recall = classification_report(y_test, y_pred_stacking,
output_dict=True)['weighted avg']['recall']
stacking_precision = classification_report(y_test, y_pred_stacking,
output_dict=True)['weighted avg']['precision']
stacking_acc = accuracy_score(y_test, y_pred_stacking)

# Print F1 score, recall, precision, and accuracy for voting and
stacking classifiers
print('F1 Score (voting): %.2f%%' % (voting_f1 * 100))
print('Recall (voting): %.2f%%' % (voting_recall * 100))
print('Precision (voting): %.2f%%' % (voting_precision * 100))
print('Accuracy (voting): %.2f%%' % (voting_acc * 100))
print("\n")
print('F1 Score (stacking): %.2f%%' % (stacking_f1 * 100))
print('Recall (stacking): %.2f%%' % (stacking_recall * 100))
print('Precision (stacking): %.2f%%' % (stacking_precision * 100))
print('Accuracy (stacking): %.2f%%' % (stacking_acc * 100))

```

Out [24]:

Classification Report (voting):

	precision	recall	f1-score	support
0	0.77	0.78	0.77	6066
1	0.77	0.77	0.77	6064
accuracy			0.77	12130
macro avg	0.77	0.77	0.77	12130
weighted avg	0.77	0.77	0.77	12130

Classification Report (stacking):

	precision	recall	f1-score	support
0	0.76	0.80	0.78	6066
1	0.79	0.75	0.77	6064
accuracy			0.77	12130
macro avg	0.77	0.77	0.77	12130
weighted avg	0.77	0.77	0.77	12130

F1 Score (voting): 77.33%
 Recall (voting): 77.33%
 Precision (voting): 77.33%
 Accuracy (voting): 77.33%

F1 Score (stacking): 77.37%
 Recall (stacking): 77.39%
 Precision (stacking): 77.45%
 Accuracy (stacking): 77.39%

In [25]:

```

import matplotlib.pyplot as plt
classifiers = ['Voting', 'Stacking']

accuracy_scores = [voting_acc, stacking_acc]
precision_scores = [voting_precision, stacking_precision]
recall_scores = [voting_recall, stacking_recall]
f1_scores = [voting_f1, stacking_f1]

fig, ax = plt.subplots()
index = list(range(len(classifiers)))
bar_width = 0.2
opacity = 0.8
plt.ylim(0.77, 0.775)
rects1 = plt.bar(index, accuracy_scores, bar_width,
alpha=opacity,
color='b',
label='Accuracy')

rects2 = plt.bar([i + bar_width for i in index], precision_scores,
bar_width,
alpha=opacity,
color='g',
label='Precision')

rects3 = plt.bar([i + bar_width*2 for i in index], recall_scores,

```

```

bar_width,
alpha=opacity,
color='r',
label='Recall')

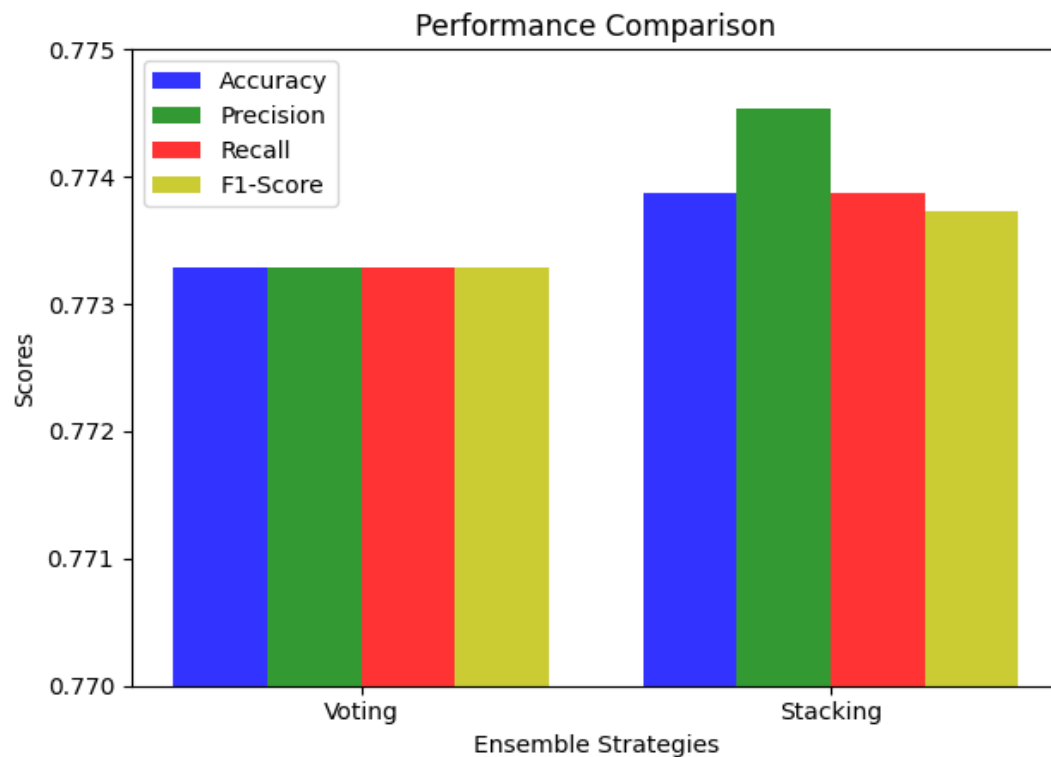
rects4 = plt.bar([i + bar_width*3 for i in index], f1_scores,
bar_width,
alpha=opacity,
color='y',
label='F1-Score')

plt.xlabel('Ensemble Strategies')
plt.ylabel('Scores')
plt.title('Performance Comparison')
plt.xticks([i + bar_width*1.5 for i in index], classifiers)
plt.legend()

plt.tight_layout()
plt.show()

```

Out [25]:



In [26]:
import scikitplot as skplt

```

# Generate precision-recall curve for the voting classifier
skplt.metrics.plot_precision_recall(y_test, voting_clf.predict_proba(X_test_scaled))
plt.title('Precision-Recall Curve (Voting)')

# Generate ROC curve for the voting classifier
skplt.metrics.plot_roc(y_test, voting_clf.predict_proba(X_test_scaled))
plt.title('ROC Curve (Voting)')

# Generate precision-recall curve for the stacking classifier
skplt.metrics.plot_precision_recall(y_test, stacking_clf.predict_proba(X_test_scaled))
plt.title('Precision-Recall Curve (Stacking)')

# Generate ROC curve for the stacking classifier
skplt.metrics.plot_roc(y_test, stacking_clf.predict_proba(X_test_scaled))
plt.title('ROC Curve (Stacking)')

```

Out [26]: Text(0.5, 1.0, 'ROC Curve (Stacking)')

